



- **Elasticache** - in memory cache in the cloud
- Improves performance of web applications, allowing you to retrieve information from fast in-memory caches rather than slower disk based databases
- Sits between your application and the database:
 - e.g. an application frequently requesting specific product information for your best selling products
- Takes the load off your databases
- Good if your database is particularly **read-heavy** and the data is not changing frequently



Elasticache Benefits and Use Cases

- Improves performance for read-heavy workloads
 - e.g. Social Networking, gaming media sharing, Q&A portals
- Frequently-accessed data is stored in memory for low-latency access, improving the overall performance of your application.
- Also good for compute heavy workloads
 - e.g. recommendation engines
- Can be used to store results of I/O intensive database queries or output of compute-intensive calculations.



2 Types of Elasticache

- **Memcached**

- Widely adopted memory object caching system
- Multi-threaded
- No Multi-AZ capability

- **Redis**

- Open-source in-memory key-value store
- Supports more complex data structures: sorted sets and lists
- Supports Master / Slave replication and Multi-AZ for cross AZ redundancy



Caching Strategies

- 2 strategies available: **Lazy Loading** and **Write-Through**:
 - **Lazy Loading** - Loads the data into the cache only when necessary
 - If requested data is in the cache, Elasticache returns the data to the application.
 - If the data is not in the cache or has expired, Elasticache returns a null.
 - Your application then fetches the data from the database and writes the data received into the cache so that it is available next time.



Lazy Loading Advantages and Disadvantages

Advantages	Disadvantages
<p>Only requested data cached: Avoids filling up cache with useless data</p>	<p>Cache miss penalty : Initial request Query to database Writing of data to the cache</p>
<p>Node failures are not fatal a new empty node will just have a lot of cache misses initially</p>	<p>Stale data - if data is only updated when there is a cache miss, it can become stale. Doesn't automatically update if the data in the database changes</p>



Lazy Loading and TTL

- **TTL (Time To Live)**
 - Specifies the number of seconds until the key (data) expires to avoid keeping stale data in the cache
 - Lazy Loading treats an expired key as a cache miss and causes the application to retrieve the data from the database and subsequently write the data into the cache with a new TTL
 - Does not eliminate stale data — but helps to avoid it



Write-Through Advantages and Disadvantages

- **Write Through** - adds or updates data to the cache whenever data is written to the database

Disdvantages	Advantages
Write penalty : Every write involves a write to the cache as well as a write to the database	Data in the cache never stale
If a node fails and a new one is spun up, data is missing until added or updated in the database (mitigate by implementing Lazy Loading in conjunction with write-through)	Users are generally more tolerant of additional latency when updating data than when retrieving it
Wasted resources if most of the data is never read	



Elasticache Exam Tips

- In-memory cache sits between your application and database
- 2 different caching strategies: Lazy loading and Write Through
- Lazy Loading only caches the data when it is requested
- Elasticache Node failures not fatal, just lots of cache misses
- Cache miss penalty: Initial request, query database, writing to cache
- Avoid stale data by implementing a TTL

Elasticache Exam Tips

- Write Through strategy writes data into the cache whenever there is a change to the database
- Data is never stale
- Write penalty: Each write involves a write to the cache
- Elasticache node failure means that data is missing until added or updated in the database
- Wasted resources if most of the data is never used